

# Slide Set 1: Introduction to OpenShift

## Slide 1 – Title Slide

- *“OpenShift Architecture: Enterprise Kubernetes for Modern Applications”*

## Slide 2 – Why OpenShift?

- Kubernetes is powerful but raw
- Missing pieces: logging, monitoring, CI/CD
- OpenShift adds enterprise-grade features
- Security-first by design

## Slide 3 – Flavors of OpenShift

- OKD (community edition)
- OpenShift Container Platform (enterprise)
- OpenShift Online (SaaS by Red Hat)
- OpenShift Dedicated (managed clusters)

### Diagram Suggestion:

A car analogy: Kubernetes = Engine, OpenShift = Full Car.

### Speaker Note:

“Think of Kubernetes as an engine. It’s powerful but you can’t drive it alone. OpenShift gives you the steering, dashboard, and safety features you need to take it on the road.”

---

## Slide Set 2: Core Architecture

## Slide 4 – Control Plane

- API Server (front door for commands)

- etcd (cluster database)
- Controller Manager (enforces desired state)
- Scheduler (decides pod placement)

### **Slide 5 – Worker Nodes**

- Kubelet (node agent)
- CRI-O (container runtime)
- Kube-proxy (networking for pods)

### **Slide 6 – OpenShift Additions**

- Operators (automate app management)
- Machine Config Operator (node config automation)

### **Diagram Suggestion:**

Two-layer diagram: Control Plane (brain) + Worker Nodes (muscle).

### **Speaker Note:**

“When you say ‘I want 5 replicas,’ the Control Plane remembers that request, and the Workers actually run it.”

---

## **Slide Set 3: Networking**

### **Slide 7 – Networking Basics**

- Pods get unique IPs
- Services provide stable DNS + IPs
- SDN ensures cluster-wide connectivity

### **Slide 8 – Routes**

- Native to OpenShift
- Map internal services to external URLs
- Supports TLS out of the box

### **Slide 9 – Example Flow**

- Frontend → Backend → Database
- Exposed to public via Route

#### **Diagram Suggestion:**

Service mesh diagram: pods inside cluster + external user accessing via Route.

#### **Speaker Note:**

“Kubernetes makes internal communication easy, but OpenShift makes external access effortless with Routes.”

---

## **Slide Set 4: Security**

### **Slide 10 – Security Enhancements**

- SCCs stricter than PodSecurityPolicies
- RBAC: fine-grained user roles
- OAuth integration: LDAP, AD, GitHub

### **Slide 11 – Image Security**

- ImageStreams
- Vulnerability scanning
- Signed images

#### **Diagram Suggestion:**

Security shield layers: User → API → SCC → Pod.

**Speaker Note:**

“OpenShift locks down workloads by default. Unlike Kubernetes, it won’t let you run as root unless you explicitly allow it.”

---

## Slide Set 5: Developer Tools

### Slide 12 – Source-to-Image (S2I)

- Build images directly from source code
- No Dockerfile needed

### Slide 13 – BuildConfigs & DeploymentConfigs

- Automate CI/CD pipelines
- Trigger redeploys on image changes

### Slide 14 – Web Console & Operators

- Developer view (drag & drop apps)
- Operators manage databases, message queues

**Diagram Suggestion:**

Flow: GitHub commit → Build → Image → Deployment → Route.

**Speaker Note:**

“A junior developer can deploy apps without touching YAML, thanks to OpenShift’s developer console and S2I.”

---

## Slide Set 6: Storage & Stateful Apps

### Slide 15 – Persistent Storage

- PVs and PVCs for storage abstraction

- StorageClasses for dynamic provisioning

### Slide 16 – StatefulSets

- Pods with persistent identity
- Essential for databases

#### Diagram Suggestion:

App pods (stateless) vs DB pods (stateful with PVC).

#### Speaker Note:

“Frontend pods can come and go, but your order database must survive. PVCs make that possible.”

---

## Slide Set 7: Monitoring & Logging

### Slide 17 – Monitoring

- Prometheus for metrics
- Grafana for dashboards
- Alertmanager for notifications

### Slide 18 – Logging

- EFK stack (Elasticsearch, Fluentd, Kibana)
- Centralized logs for troubleshooting

#### Diagram Suggestion:

Log + metric pipeline diagram (Pods → Fluentd → Elasticsearch → Kibana).

#### Speaker Note:

“With EFK, even if a pod dies, logs remain. This is crucial for debugging distributed apps.”

---

## Slide Set 8: Comparisons

### Slide 19 – OpenShift vs Kubernetes

- Kubernetes = DIY platform
- OpenShift = Enterprise-ready

### Slide 20 – OpenShift vs Others

- Rancher → multi-cluster management
- Tanzu → VMware focus
- Cloud services → not hybrid-friendly

#### Diagram Suggestion:

Comparison table (Features: Security, Dev Tools, Hybrid, Support).

#### Speaker Note:

“Kubernetes is like IKEA furniture — you assemble it. OpenShift is like a ready apartment.”

---

## Slide Set 9: Case Studies

### Slide 21 – Banking

- PCI-DSS compliance, secure workloads

### Slide 22 – Retail

- Walmart’s Black Friday scaling

### Slide 23 – Telecom

- 5G workloads on bare metal

### Slide 24 – Healthcare

- HIPAA compliance with audit logging

**Diagram Suggestion:**

Industry logos + OpenShift logo in center.

**Speaker Note:**

“These are not toy use cases. Fortune 500 companies rely on OpenShift for mission-critical apps.”

---

## **Slide Set 10: Wrap-Up**

### **Slide 25 – Summary**

- OpenShift = Kubernetes + Enterprise Tools
- Security, Networking, Storage, Monitoring
- Developer-friendly with Operators & S2I

### **Slide 26 – Discussion Questions**

- Why do enterprises pay for OpenShift?
- What challenges have you faced in Kubernetes?
- Can OpenShift replace traditional PaaS?

**Speaker Note:**

“Remember: Kubernetes solves orchestration. OpenShift solves enterprise adoption.”